

手写数字识别

1.4 LeNet-5

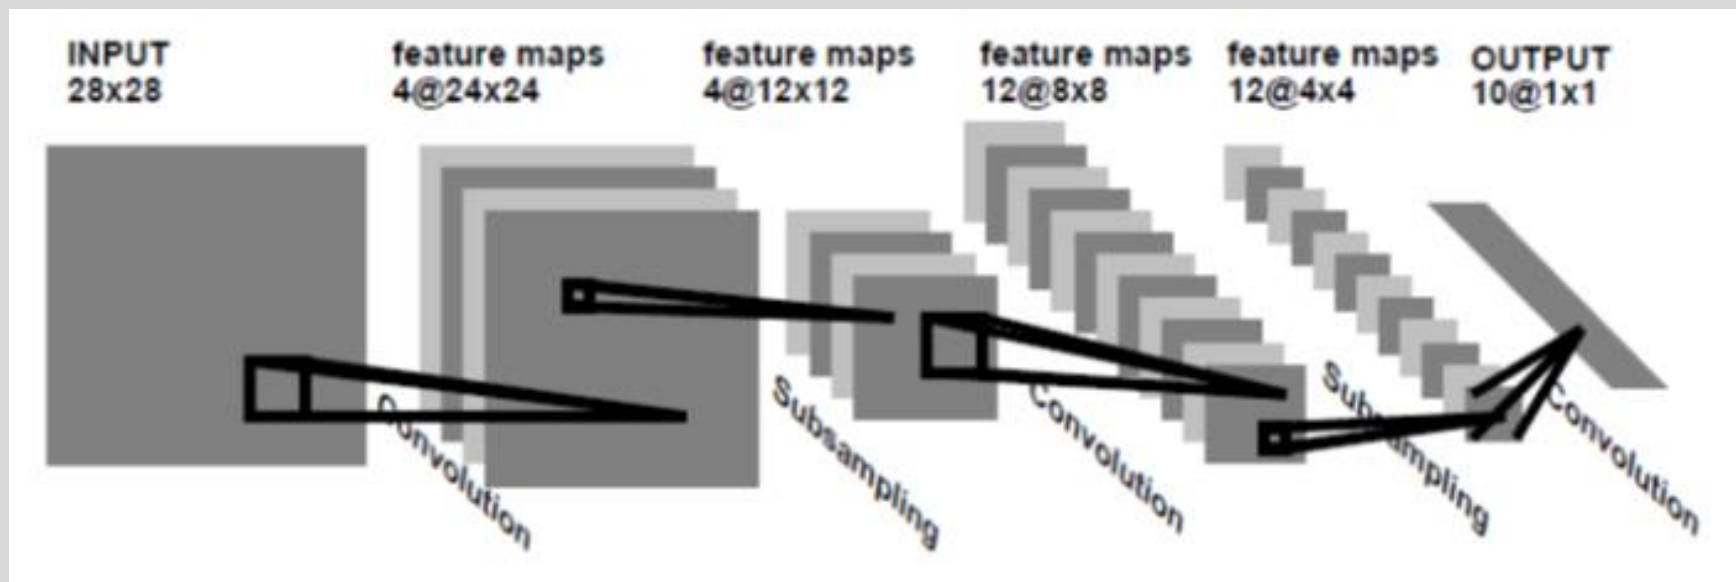
@tm9161

L5

LeNet-5

1. LeNet-5网络结构
2. LeNet-5搭建与训练
3. LeNet-5保存、调用和预测

LeNet-5 发展历史



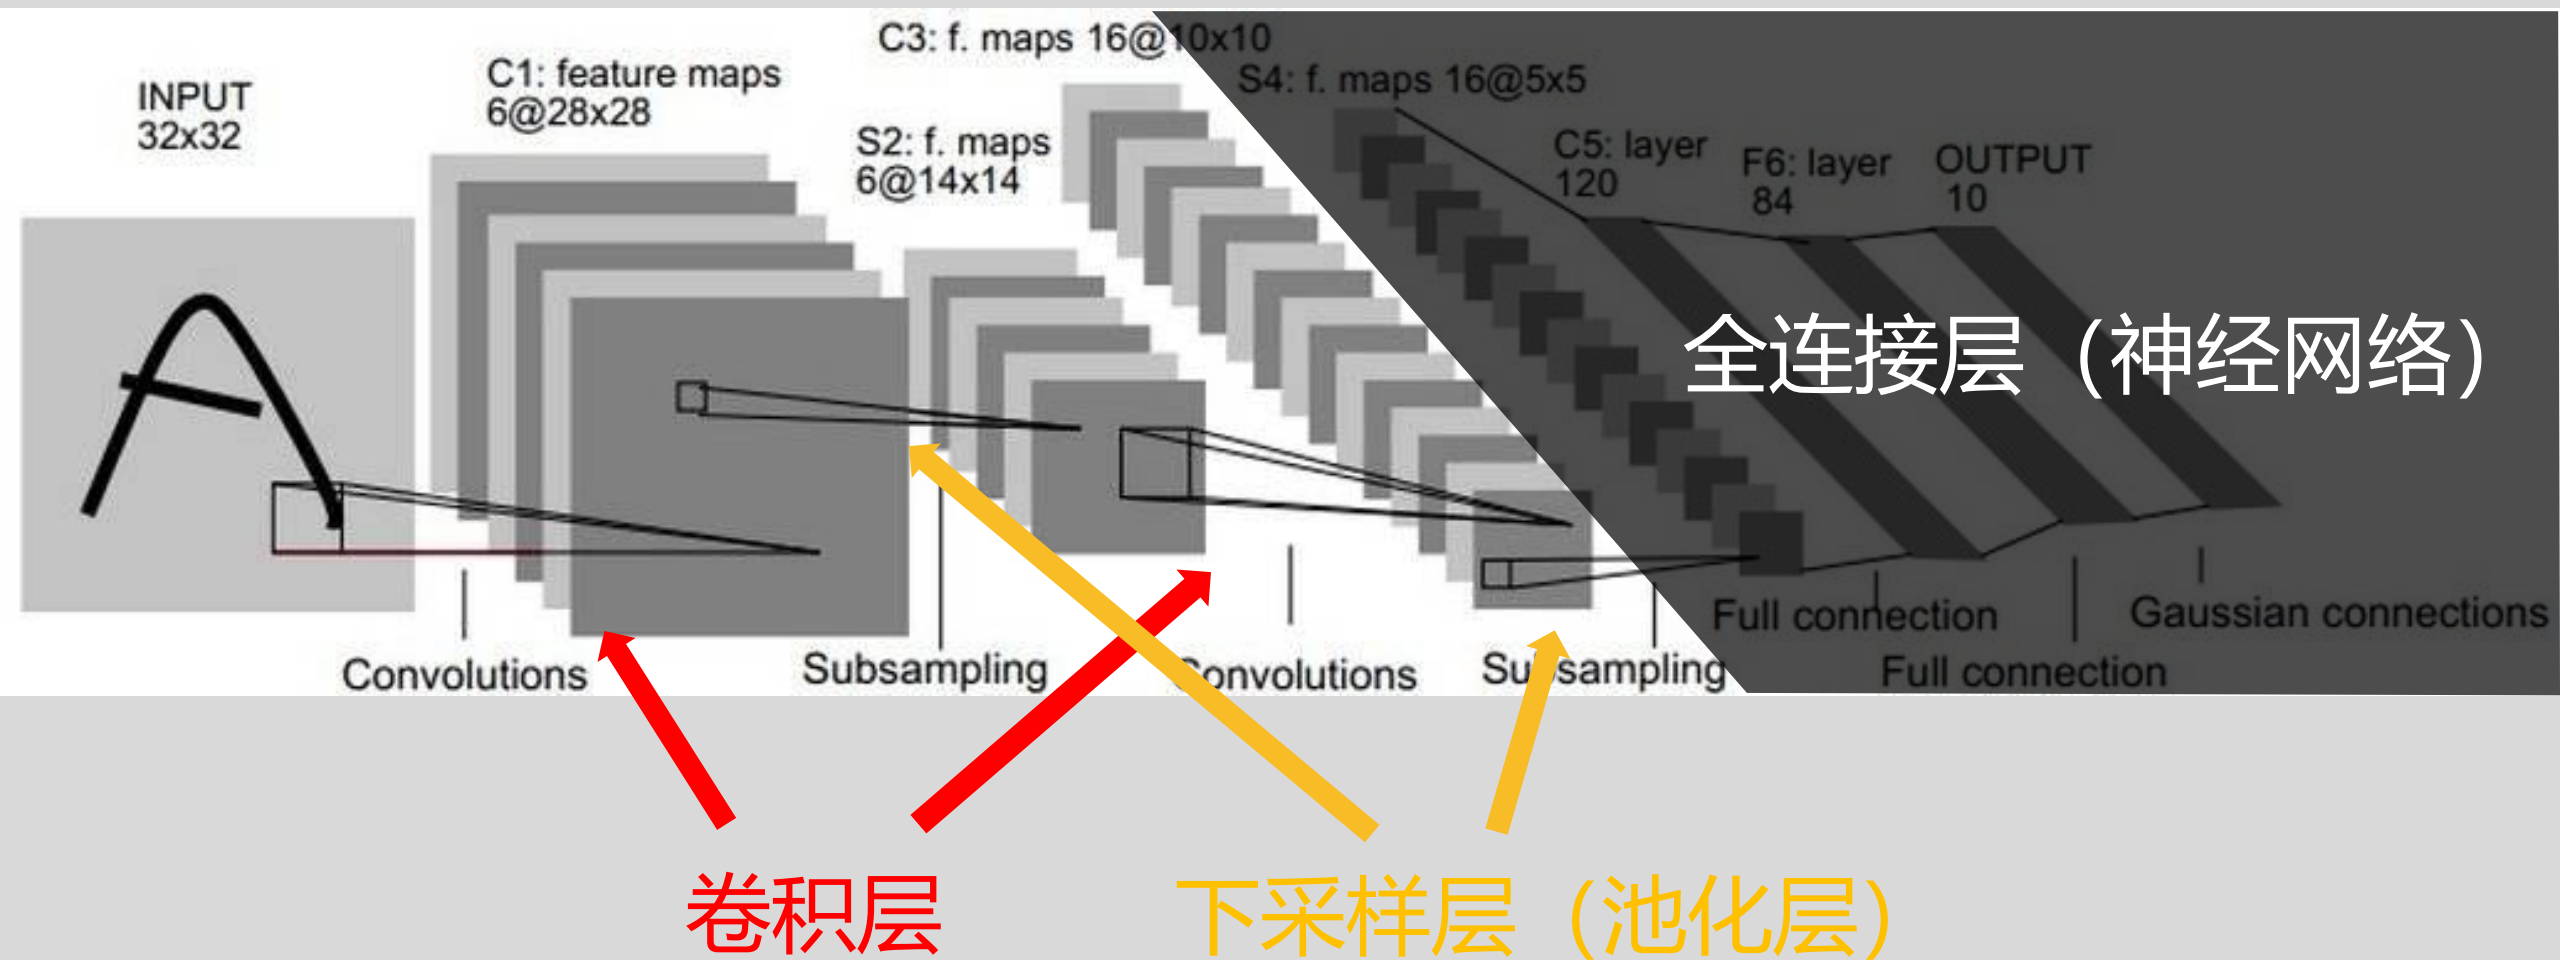
LeNet-1

1989年：Yann LeCun等人，结合反向传播算法的卷积神经网络来识别手写数字，并成功地用于识别手写邮政编码。

1990年：他们的模型在美国邮政总局提供的邮政编码数字数据的测试结果表明，错误率仅为1%，拒绝率约为9%。

1998年：他们将手写数字识别的各种方法在标准的手写数字识别基准上进行比较，结果表明他们的网络优于所有其他模型，经过多年的研究和迭代，最终发展成为LeNet-5。

网络结构



Convolution: Trying every possible match

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

Pooling


0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

max pooling

1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77


Deep stacking

Layers can be repeated several (or many) times.



Fully connected layer

Vote depends on how strongly a value predicts X or O



Draw your number here



Downsampled drawing:



First guess:


0

Second guess:

8

Layer visibility

Input layer	Show
Convolution layer 1	Show
Downsampling layer 1	Show
Convolution layer 2	Show
Downsampling layer 2	Show
Fully-connected layer 1	Show
Fully-connected layer 2	Show
Output layer	Show



Made by Adam Harley. [Project details.](#)

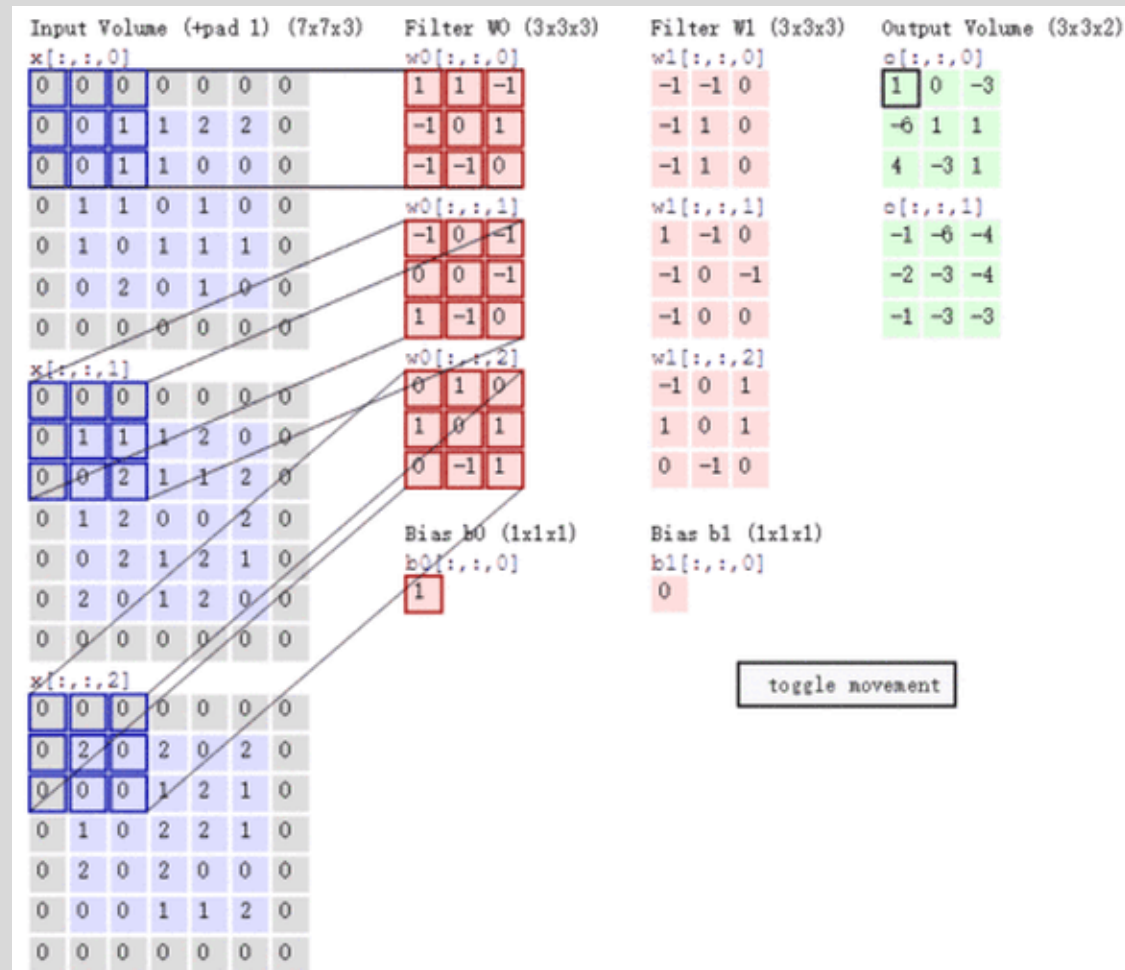
<https://www.cs.ryerson.ca/~aharley/vis/conv/>


训练参数

卷积层：卷积核中的参数+偏置项；
(3层*3x3大小 + 1*偏置) *2个

池化层：无参数需要训练。

*卷积核维度&卷积核个数区分





	输入	卷积、池化、神经元	输出	训练参数
输入层*			32*32	0
卷积层1	32*32	6个 5*5卷积核 步长为1	6*28*28 (32-5+0) /1+1	1* (5*5) *6 +6=156
池化层1	6*28*28	2*2 步长为2	6*14*14	0
卷积层2	6*14*14	16个 5*5卷积核 步长为1	16*10*10 (14-5+0) /1+1	6* (5*5) *16 +16=2416
池化层2	16*10*10	2*2 步长为2	16*5*5	0
全连接层1	16*5*5	120个 5*5卷积核 步长为1	120*1*1 (5-5+0) /1+1	16* (5*5) *120+120=48120
全连接层2	120		84	120*84+84=10164
输出层	84		10	84*10+10=850

步长 Stride & 加边 Padding

卷积后尺寸=

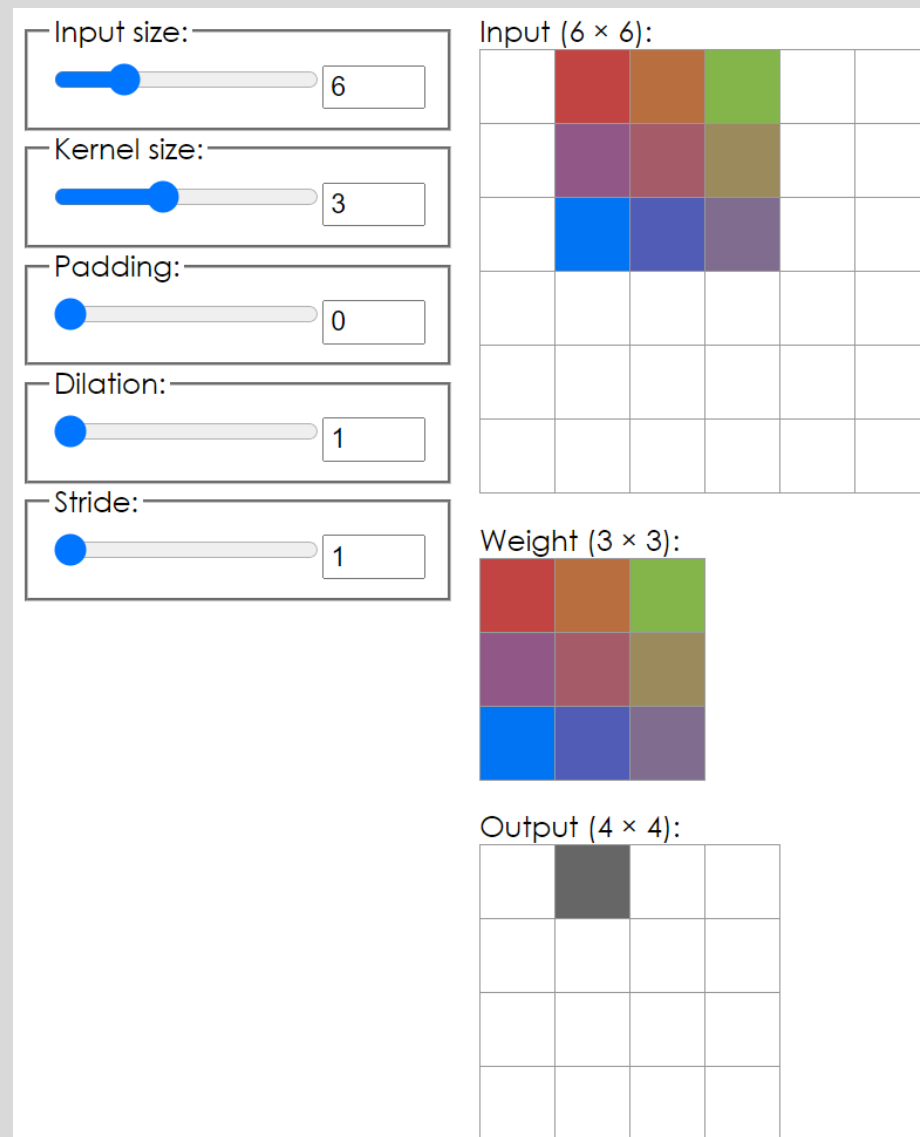
$$(\text{输入} - \text{卷积核} + \text{加边像素数}) / \text{步长} + 1$$
$$(6 - 3 + 0) / 1 + 1 = 4$$

Tensorflow 默认:

Padding= 'valid' (丢弃), strides=1

*长宽的改变

<https://ezyang.github.io/convolution-visualizer/index.html>




图片读取&预处理

```
In [3]: img = cv2.imread('3.png', 0)  
#读取图片
```

```
In [4]: plt.imshow(img)
```

```
Out[4]: <matplotlib.image.AxesImage at 0x1ca56b9ae80>
```



```
In [5]: img = cv2.resize(img, (28, 28))  
img = img.reshape(1, 28, 28, 1)  
img = img/255
```

1.图片读取: cv2.imread

2.图片大小调整: cv2.resize

3.图片维度调整: reshape

```
train_images.shape
```

```
(60000, 28, 28)
```


```
train_images = train_images.reshape(60000, 28, 28, 1)  
test_images = test_images.reshape(10000, 28, 28, 1)
```

```
train_images.shape
```


```
(60000, 28, 28, 1)
```

4.归一化: /255

维度改变 reshape



9*9*1



7*7*3

模型载入&预测

```
In [ ]: # 保存模型  
model.save('mnist.h5')
```

```
In [2]: new_model = tf.keras.models.load_model('mnist.h5')  
#调用模型
```

```
In [6]: predict = new_model.predict(img)
```


1

数据采集



```
In [1]: import tensorflow as tf

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
```

```
In [12]: (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
```

```
In [13]: train_images = train_images.reshape(60000, 28, 28, 1)
test_images = test_images.reshape(10000, 28, 28, 1)
#增加维度, 用于卷积操作
```

```
In [14]: train_images = train_images / 255
test_images = test_images / 255
```

```
In [15]: train_labels = np.array(pd.get_dummies(train_labels))
test_labels = np.array(pd.get_dummies(test_labels))
```

2

建立模型



```
In [16]: model = tf.keras.Sequential()
```

```
In [17]: model.add(tf.keras.layers.Conv2D(filters = 6, kernel_size = (5, 5), input_shape=(28, 28, 1), padding = 'same', activation = "sigmoid"))
model.add(tf.keras.layers.AveragePooling2D(pool_size = (2, 2)))
model.add(tf.keras.layers.Conv2D(filters = 16, kernel_size = (5, 5), activation = "sigmoid"))
model.add(tf.keras.layers.AveragePooling2D(pool_size = (2, 2)))
model.add(tf.keras.layers.Conv2D(filters = 120, kernel_size = (5, 5), activation = "sigmoid"))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(84, activation='sigmoid'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))
```

```
In [18]: model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 28, 28, 6)	156
average_pooling2d_6 (Average)	(None, 14, 14, 6)	0
conv2d_10 (Conv2D)	(None, 10, 10, 16)	2416
average_pooling2d_7 (Average)	(None, 5, 5, 16)	0
conv2d_11 (Conv2D)	(None, 1, 1, 120)	48120
flatten_2 (Flatten)	(None, 120)	0
dense_9 (Dense)	(None, 84)	10164
dense_10 (Dense)	(None, 10)	850

Total params: 61,706

Trainable params: 61,706

Non-trainable params: 0

3

模型训练



```
In [19]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc'])
```

```
In [20]: history = model.fit(train_images, train_labels, epochs = 10, validation_data=(test_images, test_labels))
```

```
Epoch 1/10
1875/1875 [=====] - 37s 19ms/step - loss: 1.4908 - acc: 0.4763 - val_loss: 0.2464 - val_acc: 0.9286
Epoch 2/10
1875/1875 [=====] - 31s 17ms/step - loss: 0.2135 - acc: 0.9357 - val_loss: 0.1389 - val_acc: 0.9558
Epoch 3/10
1875/1875 [=====] - 31s 17ms/step - loss: 0.1353 - acc: 0.9581 - val_loss: 0.0920 - val_acc: 0.9710
Epoch 4/10
1875/1875 [=====] - 30s 16ms/step - loss: 0.0919 - acc: 0.9715 - val_loss: 0.0752 - val_acc: 0.9771
Epoch 5/10
1875/1875 [=====] - 28s 15ms/step - loss: 0.0774 - acc: 0.9759 - val_loss: 0.0612 - val_acc: 0.9813
Epoch 6/10
1875/1875 [=====] - 28s 15ms/step - loss: 0.0611 - acc: 0.9820 - val_loss: 0.0628 - val_acc: 0.9799
Epoch 7/10
1875/1875 [=====] - 28s 15ms/step - loss: 0.0533 - acc: 0.9840 - val_loss: 0.0476 - val_acc: 0.9865
Epoch 8/10
1875/1875 [=====] - 29s 15ms/step - loss: 0.0469 - acc: 0.9856 - val_loss: 0.0488 - val_acc: 0.9840
Epoch 9/10
1875/1875 [=====] - 31s 16ms/step - loss: 0.0461 - acc: 0.9847 - val_loss: 0.0459 - val_acc: 0.9842
Epoch 10/10
1875/1875 [=====] - 28s 15ms/step - loss: 0.0366 - acc: 0.9879 - val_loss: 0.0457 - val_acc: 0.9840
```

4

模型测试




```
In [2]: new_model = tf.keras.models.load_model('mnist.h5')  
#调用模型
```

```
In [3]: img = cv2.imread('3.png', 0)  
#读取图片
```

```
In [4]: plt.imshow(img)
```

```
Out[4]: <matplotlib.image.AxesImage at 0x1ca56b9ae80>
```



```
In [5]: img = cv2.resize(img, (28, 28))  
img = img.reshape(1, 28, 28, 1)  
img = img/255
```

```
In [6]: predict = new_model.predict(img)
```

```
In [7]: predict
```

```
Out[7]: array([[2.9374178e-07, 4.0664068e-05, 4.3683460e-05, 9.9903995e-01,  
5.0486175e-08, 1.3586319e-04, 1.6269293e-09, 6.7247183e-04,  
3.3387743e-05, 3.3583085e-05]], dtype=float32)
```

```
In [8]: np.argmax(predict)
```

```
Out[8]: 3
```

参考资料：

1. Gradient-Based Learning Applied to Document Recognition

<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

2. How Deep Neural Networks Work

<https://end-to-end-machine-learning.teachable.com/p/how-deep-neural-networks-work>

3. 卷积神经网络3D可视化

<https://www.cs.ryerson.ca/~aharley/vis/conv/>

4.保存和恢复模型

https://tensorflow.google.cn/tutorials/keras/save_and_load?hl=zh-cn